

Project Report

1. Road Condition Classification Using CNN

Introduction:

The rapid increase in road usage has resulted in significant wear and tear, necessitating timely and effective maintenance strategies. This project focuses on classifying road conditions into four categories: Good, Poor, Satisfactory, and Very Poor, leveraging a Convolutional Neural Network (CNN)-based deep learning model for accurate assessment.

Objective:

- The primary goal of this project is to develop a CNN-based classification model that can accurately categorize road damage from images.
- This classification helps authorities prioritize road maintenance and ensure road safety.

Dataset:

- It contains images categorized into four classes: Good, Poor, Satisfactory, and Very Poor
- Training and testing sets were stored in separate directories: Training and Testing

Dataset Source: <https://www.kaggle.com/datasets/prudhvignv/road-damage-classification-and-assessment>

Implementation Steps:

Step 1: Data Preprocessing and Augmentation:

To enhance the robustness of the model, data augmentation techniques were applied using TensorFlow's ImageDataGenerator. Augmentations included:

- Rescaling images to a range of [0,1].
- Random shear transformation.
- Zoom transformations.
- Horizontal flipping.
- Resizing all images to **64x64 pixels**.

Step 2: Model Architecture The CNN model was built using TensorFlow and Keras with the following architecture:

- **Input Layer:** Image of size (64, 64, 3)
- **Convolutional Layers:**
 - 32 filters, kernel size (3,3), ReLU activation, followed by MaxPooling (2,2)
 - 64 filters, kernel size (3,3), ReLU activation, followed by MaxPooling (2,2)
 - 128 filters, kernel size (3,3), ReLU activation, followed by MaxPooling (2,2)
- **Fully Connected Layers:**

- Dense layer with 64 neurons (ReLU activation)
- Dense layer with 32 neurons (ReLU activation)
- Output layer with 4 neurons (Softmax activation)

Step 3: Model Compilation and Training

- The model was compiled using:
 - **Optimizer:** Adam
 - **Loss Function:** Sparse Categorical Crossentropy
 - **Metric:** Accuracy
- **Training Parameters:**
 - Batch size: 32
 - Number of epochs: 30
 - Validation on the test dataset

Results:

```

model = load_model(r"D:\svasti vector solution\CNN_RoadDamageClassification\road_damage_model.h5")

img_path = r"D:\svasti vector solution\CNN_RoadDamageClassification\sih_road_dataset\test_img_vp.jpg"

test_image = image.load_img(img_path, target_size=(64, 64))
test_image = image.img_to_array(test_image)
test_image = np.expand_dims(test_image, axis=0)
test_image = test_image / 255.0

prediction = model.predict(test_image)
predicted_class = np.argmax(prediction, axis=1)

class_labels = {0: 'Good', 1: 'Poor', 2: 'Satisfactory', 3: 'Very Poor'}
predicted_label = class_labels[predicted_class[0]]

print(f"The predicted class is: {predicted_label}")

[25]
... WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be built. 'model.compile_metrics' will be empty until you train or evaluate the model.
1/1 ----- 0s 178ms/step
The predicted class is: Very Poor

```

```

model = load_model(r"D:\svasti vector solution\CNN_RoadDamageClassification\road_damage_model.h5")

img_path = r"D:\svasti vector solution\CNN_RoadDamageClassification\sih_road_dataset\test_img_go.jpg"

test_image = image.load_img(img_path, target_size=(64, 64))
test_image = image.img_to_array(test_image)
test_image = np.expand_dims(test_image, axis=0)
test_image = test_image / 255.0

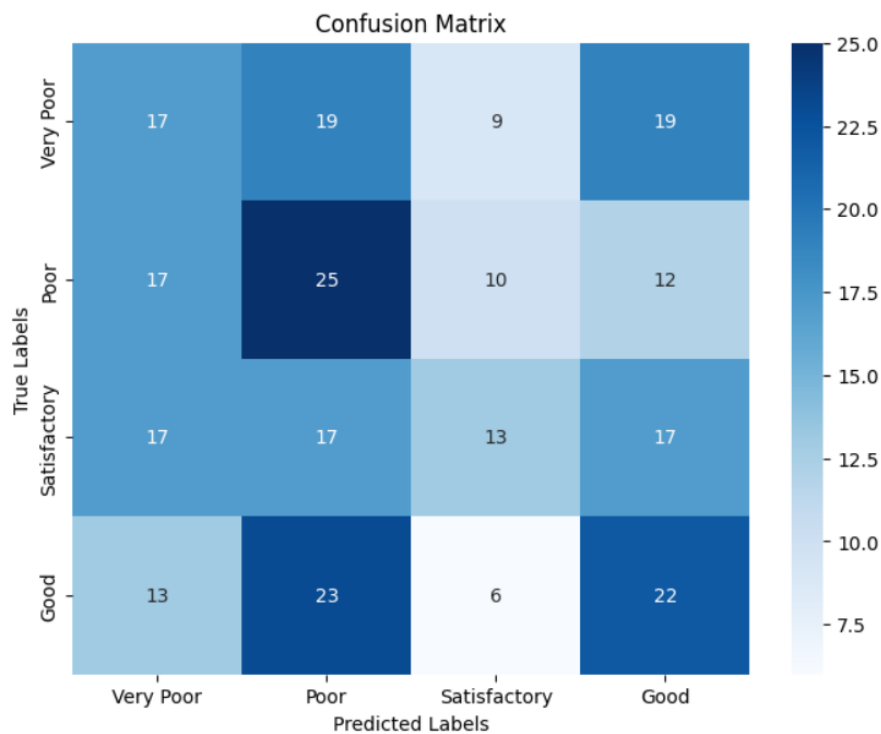
prediction = model.predict(test_image)
predicted_class = np.argmax(prediction, axis=1)

class_labels = {0: 'Good', 1: 'Poor', 2: 'Satisfactory', 3: 'Very Poor'}
predicted_label = class_labels[predicted_class[0]]

print(f"The predicted class is: {predicted_label}")

WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be built. 'model.compile_metrics' will be empty until you train or evaluate the model.
1/1 ----- 0s 136ms/step
The predicted class is: Good

```



Confusion Matrix Analysis:

- A heatmap was generated to visualize classification performance.
- Some misclassifications were observed, indicating potential improvement areas.

Model Deployment and Testing:

- The trained model was saved as road_damage_model.h5.
- The model was tested on a sample image using the following steps:
 - Image was preprocessed and converted into an array.
 - Model prediction was obtained using predict().
 - The predicted class label was mapped using a dictionary.
 - The result was displayed with Matplotlib.

Code:

https://drive.google.com/file/d/1jyVxMHCFwm2GDjPIGnLOSvXXi1bNqWwV/view?usp=drive_link

Conclusion and Future Work:

This project successfully implemented a CNN-based road damage classification system. The model demonstrated promising accuracy and can be further enhanced by:

- Using a larger dataset to improve generalization.
- Fine-tuning hyperparameters to optimize performance.
- Implementing transfer learning techniques for better feature extraction.
- Deploying the model in real-time applications for automated road condition monitoring.

2.Road Condition Classification using CNN and MobileNetV2

Introduction:

Road damage and poor road conditions pose significant challenges for transportation safety and maintenance planning. This project aims to classify road conditions into four categories: Good, Poor, Satisfactory, and Very Poor using deep learning techniques. By leveraging MobileNetV2 with transfer learning, we aim to build an efficient and accurate model for road condition classification.

Objectives:

- To develop an image classification model that categorizes road conditions into four classes.
- To utilize transfer learning with MobileNetV2 for improved accuracy and computational efficiency.
- To apply image augmentation techniques to enhance model generalization.
- To evaluate model performance and improve accuracy for real-world deployment.

Step 1: Data Collection

The dataset consists of images categorized into four classes:

- **Good:** Roads in excellent condition with minimal or no visible damage.
- **Poor:** Roads with significant cracks or wear but still usable.
- **Satisfactory:** Roads with minor cracks or irregularities.
- **Very Poor:** Roads with severe damage, potholes, or hazardous conditions.

Dataset Source: <https://www.kaggle.com/datasets/prudhvignv/road-damage-classification-and-assessment>

Step 2: Data Preprocessing

- Images were resized to **224x224** pixels to match MobileNetV2's input dimensions.
- Data augmentation was applied to enhance model robustness:
 - **Rotation (15 degrees)**
 - **Width and height shift (10%)**
 - **Shear transformation (10%)**
 - **Zoom range (10%)**
 - **Horizontal flipping**
- The dataset was split into **training** and **testing** sets.

Step 3: Model Architecture

Pre-trained Model: MobileNetV2

- MobileNetV2 was selected due to its efficiency in handling image classification tasks with minimal computational overhead.
- The model was used without its top layers to allow for custom classification layers.
- The convolutional layers were **frozen** to retain pre-trained features.

Custom Layers

- **Global Average Pooling** to reduce feature dimensions.
- **Fully Connected Layer (128 neurons, ReLU activation)** for learning complex patterns.
- **Dropout Layer (50%)** to prevent overfitting.
- **Output Layer (4 neurons, Softmax activation)** to classify images into four categories.

Step 4: Model Training

Training Configuration

- **Optimizer:** Adam with a learning rate of **0.001**
- **Loss Function:** Categorical Crossentropy
- **Batch Size:** 16
- **Number of Epochs:** 10

Training Process

- ImageDataGenerator was used for real-time image augmentation.
- Training images were fed in batches to improve computational efficiency.
- The model was validated using the test dataset after each epoch.

Step 5: Model Evaluation

Performance Metrics

The model was evaluated using:

- **Categorical Accuracy:** Measures the percentage of correctly classified images.
- **Loss Function:** Evaluates how well the model is minimizing errors.

Step 6: Results

```

img_path = 'test_img_Go.jpg'
img = image.load_img(img_path, target_size=(224, 224)) # Use the same target size as used in training
img_array = image.img_to_array(img) # Convert the image to a NumPy array
img_array = np.expand_dims(img_array, axis=0) # Add batch dimension
img_array = img_array / 255.0 # Normalize if required (based on your training preprocessing)

# Use the loaded model to predict the class of the new image
prediction = model.predict(img_array)

# If you have multiple classes, you can find the class with the highest probability
predicted_class = np.argmax(prediction, axis=1)

if predicted_class[0] == 0:
    m = "Good"
elif predicted_class[0] == 1:
    m = "Poor"
elif predicted_class[0] == 2:
    m = "Moderate"
elif predicted_class[0] == 3:
    m = "Very Poor"
# Print the predicted class
print("Predicted Class:", predicted_class, " which is", m)

```

1/1 [=====] - 0s 59ms/step
Predicted Class: [0] which is Good

```

img_path = 'test_img_P.jpg'
img = image.load_img(img_path, target_size=(224, 224)) # Use the same target size as used in training
img_array = image.img_to_array(img) # Convert the image to a NumPy array
img_array = np.expand_dims(img_array, axis=0) # Add batch dimension
img_array = img_array / 255.0 # Normalize if required (based on your training preprocessing)

# Use the loaded model to predict the class of the new image
prediction = model.predict(img_array)

# If you have multiple classes, you can find the class with the highest probability
predicted_class = np.argmax(prediction, axis=1)

if predicted_class[0] == 0:
    m = "Good"
elif predicted_class[0] == 1:
    m = "Poor"
elif predicted_class[0] == 2:
    m = "Moderate"
elif predicted_class[0] == 3:
    m = "Very Poor"
# Print the predicted class
print("Predicted Class:", predicted_class, " which is", m)

```

1/1 [=====] - 0s 81ms/step
Predicted Class: [3] which is Very Poor

Step 7: Model Deployment

- The trained model was saved in HDF5 format (road_condition_model2.h5).
- It can be deployed for real-world applications such as automated road condition monitoring systems.

Code:

https://drive.google.com/file/d/1QCOxkLtac6hE4-oQYTvhl5Rg1j81nyoS/view?usp=drive_link

Load Model:

https://drive.google.com/file/d/1dxRBiOY_uyltQZ3dBQvYXYvGjUc5cSSg/view?usp=drive_link

H5 File:

https://drive.google.com/file/d/1R6UF_QEaeUrb2ML64k7za99C-lf54Mt8/view?usp=drive_link

Future Enhancements:

- **Fine-tuning MobileNetV2 layers** to improve feature extraction.
- **Increasing dataset size** with real-world road images to improve generalization.
- **Experimenting with alternative architectures** like ResNet50 or EfficientNet.
- **Deploying the model as a web application** for real-time road condition analysis.

Conclusion:

This project successfully developed a road condition classification model using MobileNetV2 with transfer learning. The model demonstrated promising results, making it a viable tool for automated road condition monitoring. Future work will focus on enhancing accuracy and deploying the model in practical applications.